

---

---

# **Collaborative Metric Learning**

**By Hsieh, Cheng-Kang, et al.**

**Presented by Kim Han**

---

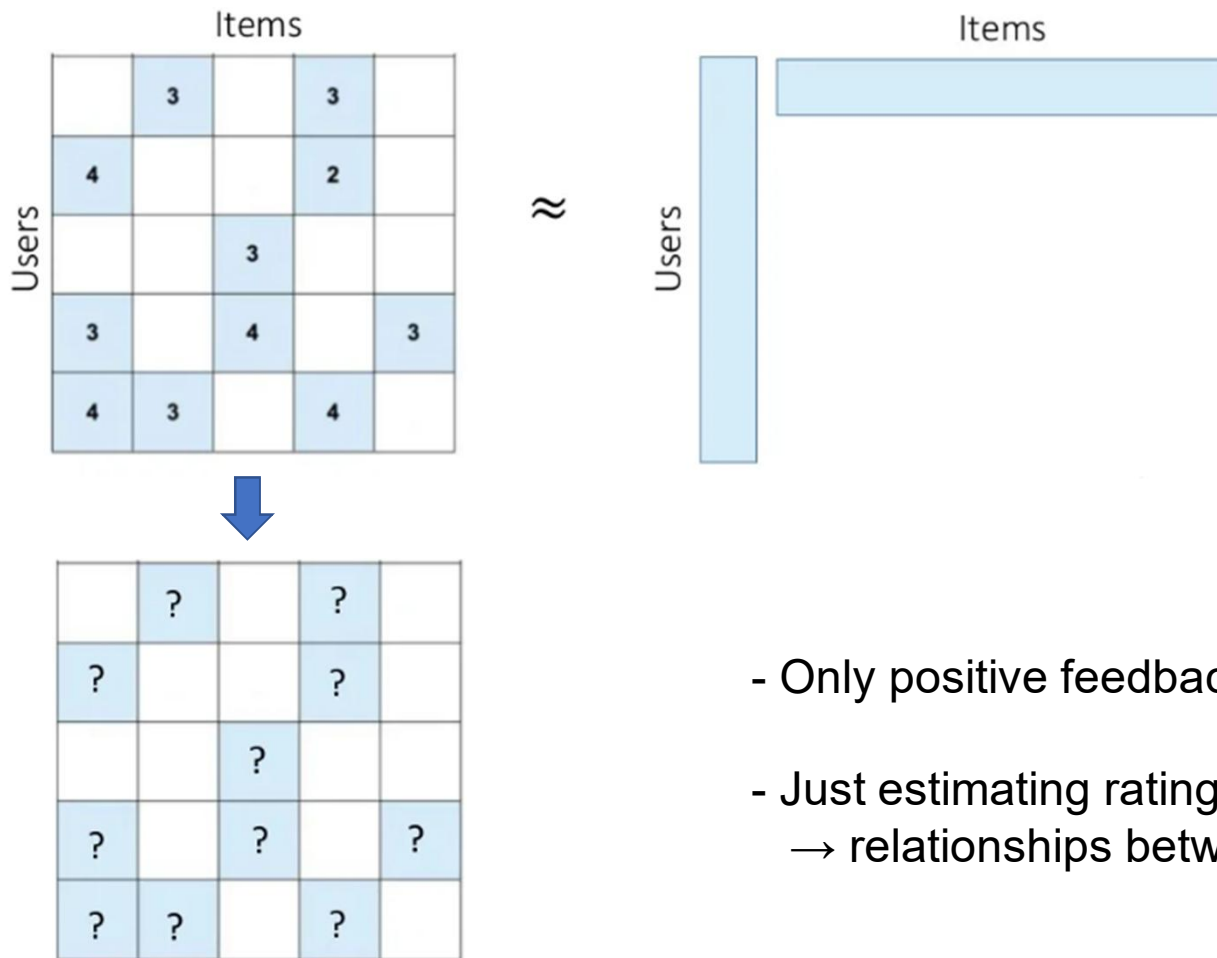
---

# Contents

---

1. Introduction
2. Background
3. Model
4. Experiment
5. Implementation
6. Conclusion

# Introduction



- Only positive feedback is available
- Just estimating ratings
  - relationships between user-user or item-item pairs

# Background

---

## Metric Learning

$$d_E(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{\sum_{i=1}^d (\mathbf{x}_{1,i} - \mathbf{x}_{2,i})^2}$$

Euclidean distance

$$d_M(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{(\mathbf{x}_1 - \mathbf{x}_2)^T S^{-1} (\mathbf{x}_1 - \mathbf{x}_2)} \quad \Rightarrow \quad d_A(x_i, x_j) = \sqrt{(x_i - x_j)^T A (x_i - x_j)}$$

Mahalanobis distance

$$\mathcal{S} = \{(x_i, x_j) | x_i \text{ and } x_j \text{ are considered similar}\}$$

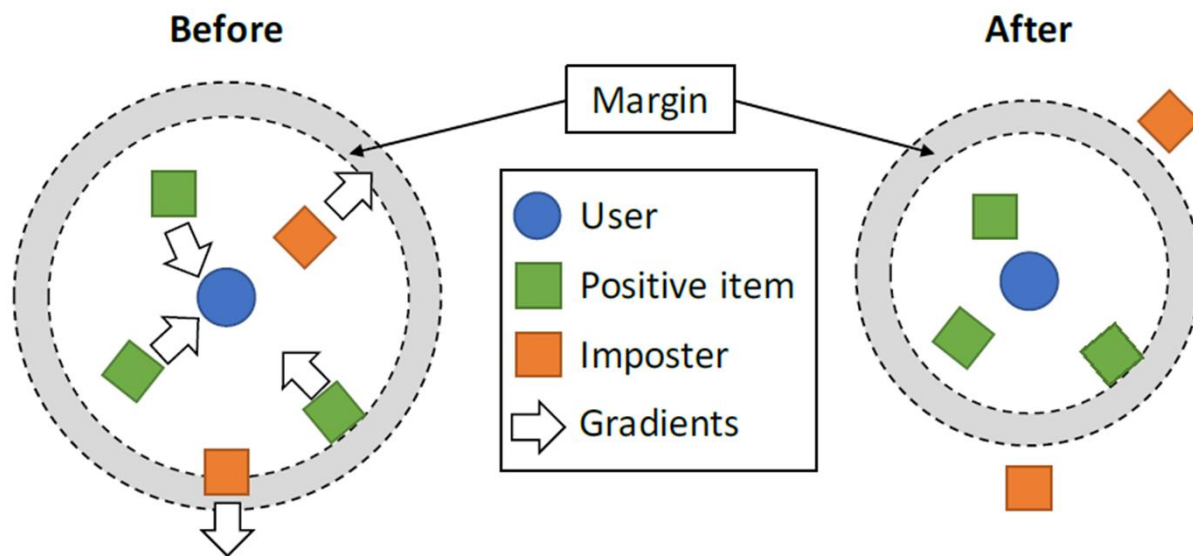
$$\mathcal{D} = \{(x_i, x_j) | x_i \text{ and } x_j \text{ are considered dissimilar}\}$$

$$\min_A \sum_{(x_i, x_j) \in \mathcal{S}} d_A(x_i, x_j)^2$$

$$\text{s.t. } \sum_{(x_i, x_j) \in \mathcal{D}} d_A(x_i, x_j)^2 \geq 1 \text{ and } A \succeq 0.$$

# Background

## Metric Learning (Large Margin Nearest Neighbor)



$$\mathcal{L}_{pull}(d) = \sum_{j \sim i} d(x_i, x_j)^2$$

$$\mathcal{L}_{push}(d) = \sum_{i, j \sim i} \sum_k (1 - y_{ik}) [1 + d(x_i, x_j)^2 - d(x_i, x_k)^2]_+$$

# Model

## Collaborative Metric Learning



1.  $D(\vec{x}_i, \vec{x}_j) + D(\vec{x}_j, \vec{x}_k) \geq D(\vec{x}_i, \vec{x}_k)$  (triangular inequality).
2.  $D(\vec{x}_i, \vec{x}_j) \geq 0$  (non-negativity).
3.  $D(\vec{x}_i, \vec{x}_j) = D(\vec{x}_j, \vec{x}_i)$  (symmetry).
4.  $D(\vec{x}_i, \vec{x}_j) = 0 \iff \vec{x}_i = \vec{x}_j$  (distinguishability).

- the items liked by this user previously
- the items liked by other users who share a similar taste with this user previously

# Model

---

## Collaborative Metric Learning

$u_i \in \mathcal{R}^r$  : User vector

$v_j \in \mathcal{R}^r$  : Item vector

$d(i, j) = \|u_i - v_j\|$  : Distance (with euclidean distance)

$$w_{ij} = \log(\text{rank}_d(i, j) + 1)$$

Approximated Ranking Weight

$$\left\lfloor \frac{J}{N} \right\rfloor \longrightarrow \left\lfloor \frac{J \times M}{U} \right\rfloor$$

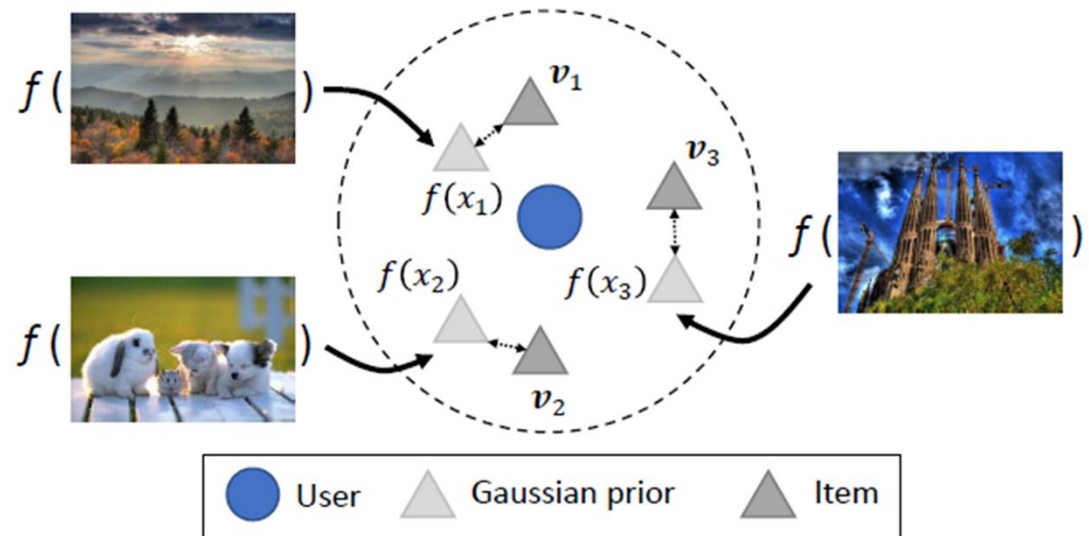
Hinge Loss

$$\mathcal{L}_m(d) = \sum_{(i,j) \in \mathcal{S}} \sum_{(i,k) \notin \mathcal{S}} w_{ij} [m + d(i, j)^2 - d(i, k)^2]_+, \quad (1)$$

# Model

## Integrating Item Features

$$\mathcal{L}_f(\theta, \mathbf{v}_*) = \sum_j \|f(\mathbf{x}_j, \theta) - \mathbf{v}_j\|^2$$





# Model

---

## Regularization

$$\|\mathbf{u}_*\|^2 \leq 1 \text{ and } \|\mathbf{v}_*\|^2 \leq 1$$

Normalize

$$C_{ij} = \frac{1}{N} \sum_n (y_i^n - \mu_i)(y_j^n - \mu_j)$$

$$\|C\|_f^2 = |a_{11}|^2 + |a_{12}|^2 + \cdots + |a_{mm}|^2$$

$$\mu_i = \frac{1}{N} \sum_n y_i^n$$

$$\|diag(C)\|_2^2 = |a_{11}|^2 + |a_{22}|^2 + |a_{33}|^2 + \cdots + |a_{mm}|^2$$

$$\mathcal{L}_c = \frac{1}{N} (\|C\|_f - \|diag(C)\|_2^2)$$

# Model

---

## Training

$$\begin{aligned} \min_{\theta, \mathbf{u}_*, \mathbf{v}_*} \quad & \mathcal{L}_m + \lambda_f \mathcal{L}_f + \lambda_c \mathcal{L}_c \\ \text{s.t.} \quad & \|\mathbf{u}_*\|^2 \leq 1 \text{ and } \|\mathbf{v}_*\|^2 \leq 1 \end{aligned}$$

1. Sample N positive pairs from S.
2. For each pair, sample U negative items and approximate  $\text{rank\_d}(i, j)$ .
3. For each pair, keep the negative item k that maximizes the hinge loss and form a mini-batch of size N.
4. Compute gradients and update parameters with a learning rate controlled by AdaGrad.
5. Censor the norm of  $\mathbf{u}^*$  and  $\mathbf{v}^*$  by  $\mathbf{y}' = \frac{\mathbf{y}}{\max(\|\mathbf{y}\|, 1)}$
6. Repeat this procedure until convergence.

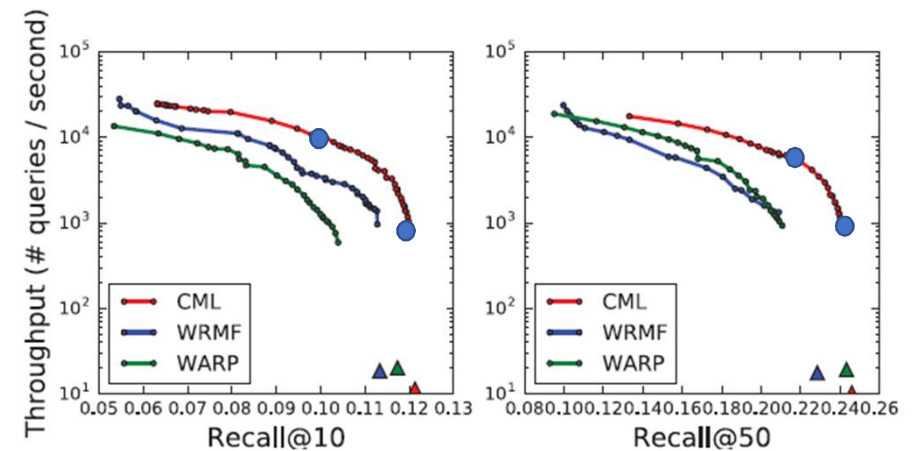
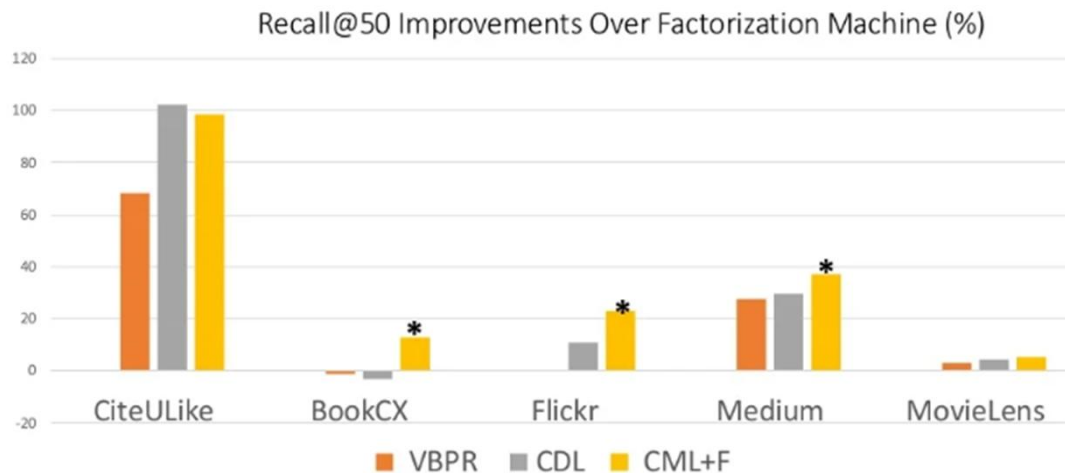
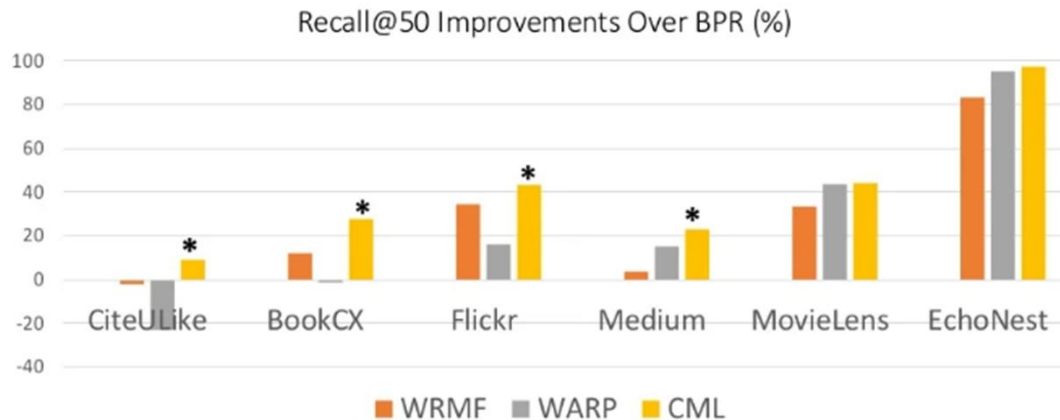
# Experiment

Table 1: Dataset Statistics.

	CiteULike	BookCX	Flickr	Medium	MovieLens20M	EchoNest
Domain	Paper	Book	Photography	News	Movie	Song
# Users	7,947	22,816	43,758	61,909	129,797	766,882
# Items	25,975	43,765	100,000	80,234	20,709	260,417
# Ratings	142,794	623,405	1,372,621	2,047,908	9,939,873	7,261,443
Concentration <sup>a</sup>	33.47%	33.10%	13.48%	55.38%	72.52%	65.88%
Features Type	Tags	Subjects	Image Features	Tags	Genres, Keywords	NA
# Feature Dim.	10,399	7,923	2,048	2,313	10,399	NA

	WRMF	BPR	WARP	CML	<i>ours vs. best</i>	FM	VBPR	CDL	CML+F	<i>ours vs. best</i>
<i>Recall@50</i>										
CiteULike	0.2437	0.2489	0.1916	<b>0.2714***</b>	9.03%	0.1668	0.2807	<b>0.3375**</b>	0.3312	-1.86%
BookCX	0.0910	0.0812	0.0801	<b>0.1037***</b>	13.95%	0.1016	0.1004	0.0984	<b>0.1147***</b>	12.89%
Flickr	0.0667	0.0496	0.0576	<b>0.0711***</b>	6.59%	NA	0.0612	0.0679	<b>0.0753***</b>	10.89%
Medium	0.1457	0.1407	0.1619	<b>0.1730***</b>	6.41%	0.1298	0.1656	0.1682	<b>0.1780***</b>	5.82%
MovieLens	0.4317	0.3236	0.4649	<b>0.4665</b>	0.34%	0.4384	0.4521	0.4573	<b>0.4617*</b>	0.96%
EchoNest	0.2285	0.1246	0.2433	<b>0.2460</b>	1.10%	NA	NA	NA	NA	NA
<i>Recall@100</i>										
CiteULike	0.3112	0.3296	0.2526	<b>0.3411***</b>	3.37%	0.2166	0.3437	0.4173	<b>0.4255**</b>	1.96%
BookCX	0.1286	0.1230	0.1227	<b>0.1436***</b>	11.66%	0.1440	0.1455	0.1428	<b>0.1712***</b>	17.66%
Flickr	0.0821	0.0790	0.0797	<b>0.0922***</b>	12.30%	NA	0.0880	0.0909	<b>0.1048***</b>	15.29%
Medium	0.2112	0.2078	0.2336	<b>0.2480***</b>	6.16%	0.1900	0.2349	0.2408	<b>0.2531***</b>	5.10%
MovieLens	0.5649	0.4455	0.5989	<b>0.6022</b>	0.55%	0.5561	0.5712	0.5943	<b>0.5976</b>	0.55%
EchoNest	0.2891	0.1655	0.3021	<b>0.3022</b>	0.00%	NA	NA	NA	NA	NA

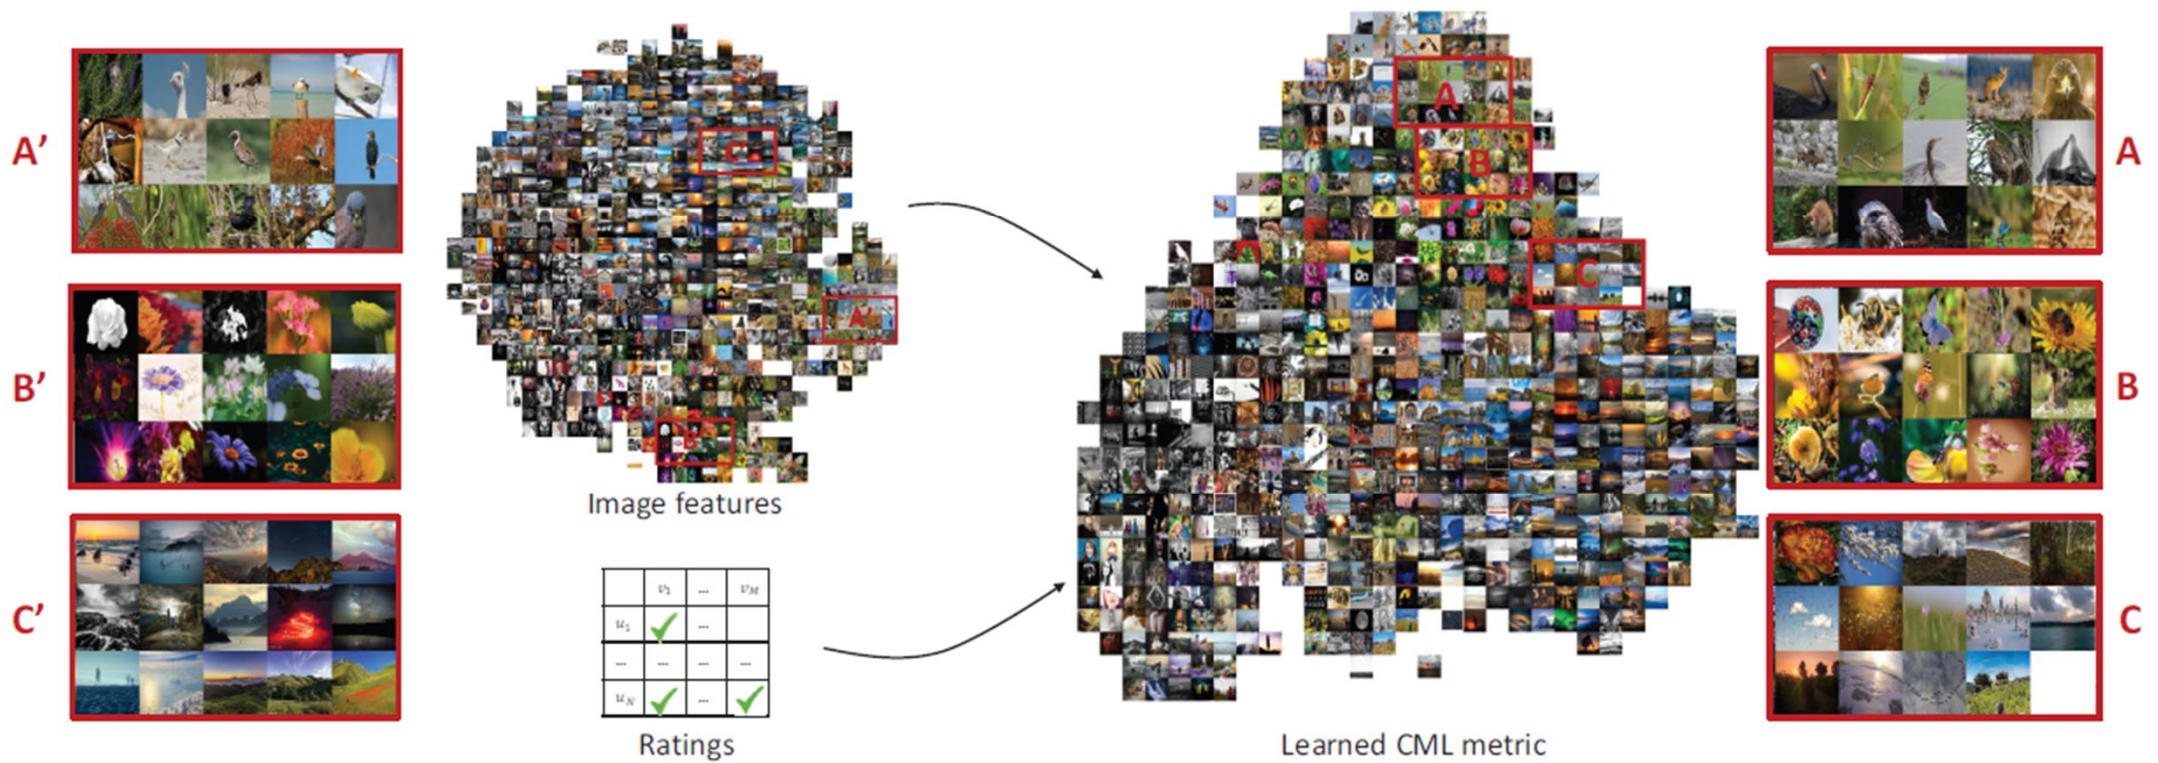
# Experiment



1. CML's throughput is improved by 106x with only 2% reduction in accuracy.
2. Over 8x, 9x faster than MF models given the same accuracy



# Experiment



# Conclusion

---

1. Implementing collaborative filtering in another way that is more efficient than the existing method
2. Shown to be very efficient and fast in the Top-K recommendation method
3. Not only user-item, but also user-user, item-item relationships were interpreted to make the entire system easier to understand

**감사합니다.**