

Neural Collaborative Filtering

Lab Intern 이준모
bubblego0217@kaist.ac.kr

Contents

1. Introduction
 2. Preliminaries
 3. Neural Collaborative Filtering
 4. Experiments
 5. Conclusion & Future Work
 6. Implementation (To be updated)
-

Introduction

- Collaborative Filtering
 - Matrix Factorization (MF)
 - Interaction : Inner product of user and item latent vectors
 - Combine with neighbor-based models, Factorization Machines
 - Performance can be hindered by the **interaction function**
 - not be sufficient to capture the complex structure of interaction
- Using Deep Neural Network
 - Little work on Recommendation (Previous)
 - DNNs to model auxiliary information (ex. Textual description of items)
 - Still resorted to MF
 - Use DNNs for **learning the interaction function** from data

Implicit Feedback & Contribution

- Implicit Feedback
 - Can be tracked automatically & Much easier to collect
 - User Satisfaction is not observed & A natural Scarcity of Negative Feedback
- Contribution
 1. **Neural Net Architecture** to model latent features & General framework NCF
 2. MF = Specialization of NCF & **Non-linearity** with Multi-Layer Perceptron
 3. **Effectiveness** of NCF Approaches & Promise of Deep Learning for CF

Learning from Implicit Data

M and N : the number of users and items

$$\mathbf{Y} \in \mathbb{R}^{M \times N} \quad \leftarrow \quad y_{ui} = \begin{cases} 1, & \text{if interaction (user } u, \text{ item } i) \text{ is observed;} \\ 0, & \text{otherwise.} \end{cases}$$

	i_1	i_2	i_3	i_4	i_5	
u_1	1	1	1	0	1	↑ users ↓
u_2	0	1	1	0	0	
u_3	0	1	1	1	0	
u_4	1	0	1	1	1	
	← items →					

(a) user-item matrix

※ Interaction doesn't mean Preference

- Observed : at least interest
- Unobserved : Missing Data & Scarcity of Negative Feedback

Learning from Implicit Data

- Recommendation problem with Implicit Feedback
 - Estimating the scores of unobserved entries in Y
 - Abstracted as learning $\hat{y}_{ui} = f(u, i | \Theta)$
- Estimate parameters Θ
 - Follow the Machine Learning Paradigm – Optimizes an objective function
 - Pointwise Loss
 - $\min \frac{1}{2} (\hat{y}_{ui} - y_{ui})^2$ (ex. OCCF)
 - Pairwise Loss
 - $\max(0, f(y_{unobs}) - f(y_{obs}) + \alpha)$ (ex. BPR)

Matrix Factorization

- Abstraction

$$\hat{y}_{ui} = f(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^T \mathbf{q}_i = \sum_{k=1}^K p_{uk} q_{ik}$$

- p_u and q_i denote latent vector for user u and item i
- **Linear model** of latent factors

- Expressiveness Limitation of MF

1. Similarity between two users by inner product

- Cosine similarity

2. Use Jaccard coefficient

- Ground-truth similarity of two users that MF needs to cover
- Defined as $s_{ij} = \frac{|\mathcal{R}_i \cap \mathcal{R}_j|}{|\mathcal{R}_i \cup \mathcal{R}_j|}$, \mathcal{R}_u be the set of items that user u has interacted with.

Matrix Factorization


1. Consider only the first three users

	i_1	i_2	i_3	i_4	i_5
u_1	1	1	1	0	1
u_2	0	1	1	0	0
u_3	0	1	1	1	0
u_4	1	0	1	1	1

↑ users

← items →

(a) user-item matrix



(b) user latent space

- Similarity

$$s_{23} \left(\frac{2}{3} = 0.66 \right) > s_{12} \left(\frac{2}{4} = 0.5 \right) > s_{13} \left(\frac{2}{5} = 0.4 \right)$$

- Interpretation

- Relative Similarity


Matrix Factorization

2. Consider new user u_4

	i_1	i_2	i_3	i_4	i_5
u_1	1	1	1	0	1
u_2	0	1	1	0	0
u_3	0	1	1	1	0
u_4	1	0	1	1	1

↑ users
← items

(a) user-item matrix



(b) user latent space

- Similarity

$$s_{41} \left(\frac{3}{5} = 0.6 \right) > s_{43} \left(\frac{2}{5} = 0.4 \right) > s_{42} \left(\frac{1}{5} = 0.2 \right)$$

- Interpretation

- p_4 closer to p_2 than p_3

Matrix Factorization

Question : Increasing the number of latent factors K ?

	i_1	i_2	i_3	i_4	i_5
u_1	1	1	1	0	1
u_2	0	1	1	0	0
u_3	0	1	1	1	0
u_4	1	0	1	1	1

↑ users
↓ users

← items →

(a) user-item matrix

Large Latent space

?

- It may hurt the generalization of the model (Overfitting)
 - Especially in sparse settings
- So use **DNNs** to learn the interaction function from data

General Framework

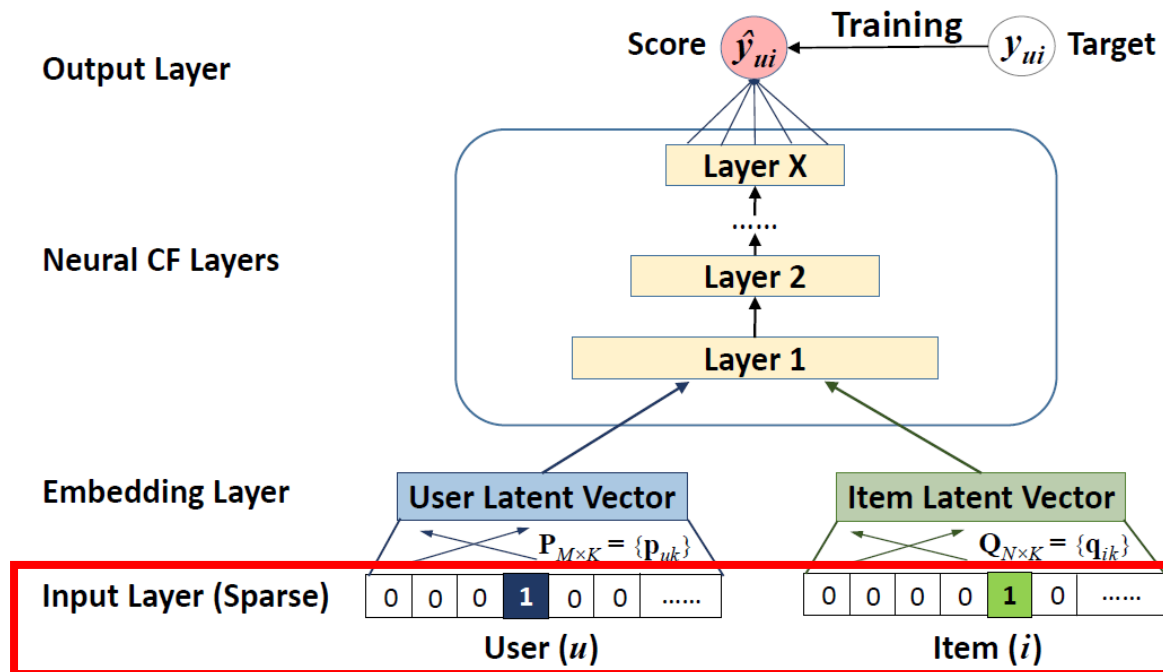


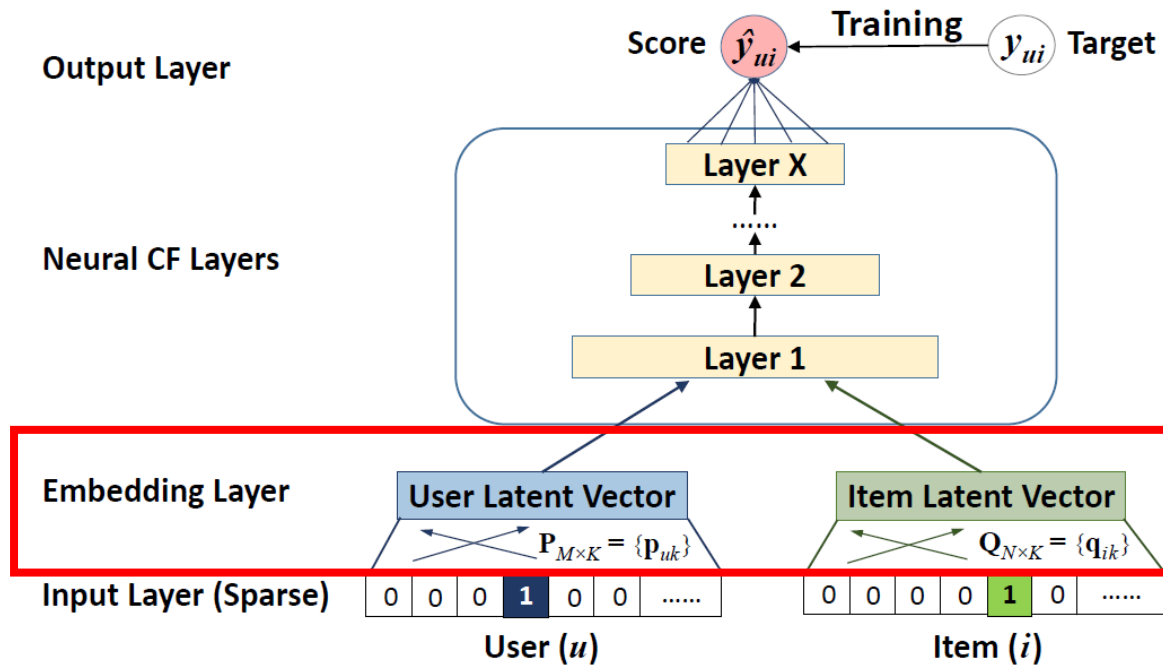
Figure 2: Neural collaborative filtering framework

- Multi-Layer Representation

1. Input Layer

- Sparse feature vectors v_u^U and v_i^I
 - Binarized by **one-hot encoding**
 - Pure collaborative filtering setting

General Framework



2. Embedding Layer

- Fully-Connected Layer
- Sparse representation to a dense vector

Figure 2: Neural collaborative filtering framework

General Framework

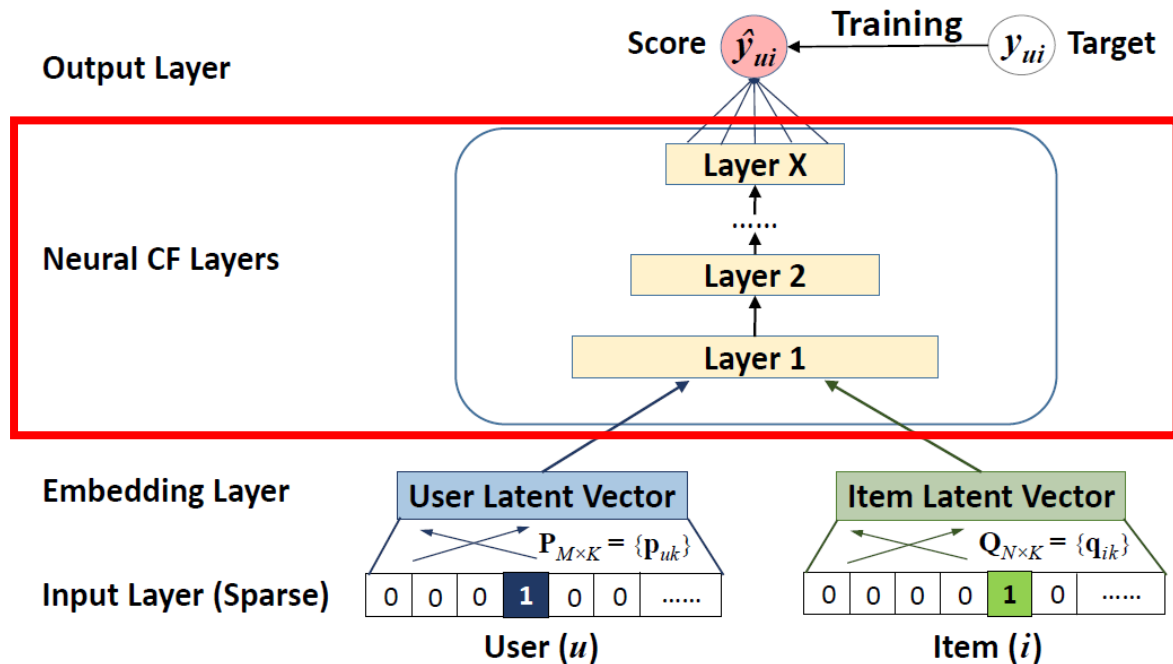


Figure 2: Neural collaborative filtering framework

3. Neural CF Layer

- Multi-Layer Neural Architecture
- Map the latent vectors to prediction scores
- Dimension of last hidden Layer X determines the model's capability

General Framework

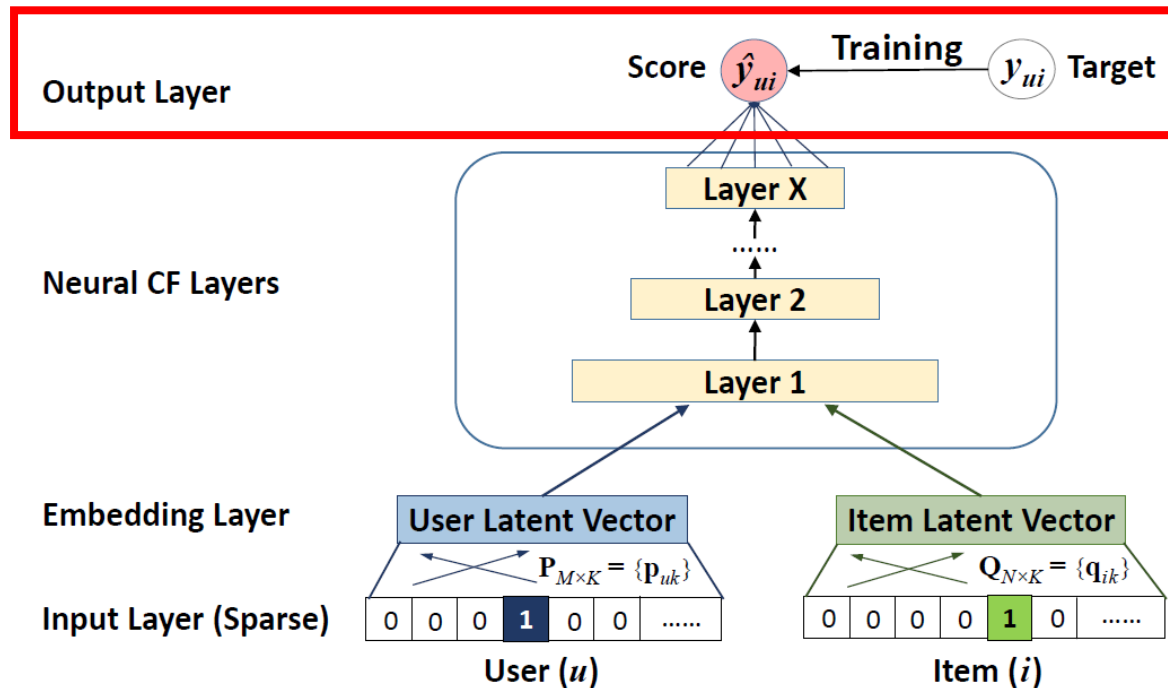


Figure 2: Neural collaborative filtering framework

4. Output Layer

- Predicted Score \hat{y}_{ui}
- By minimizing pointwise loss
 - ※ Future work for pairwise loss

• Abstraction

$$\hat{y}_{ui} = f(\mathbf{P}^T \mathbf{v}_u^U, \mathbf{Q}^T \mathbf{v}_i^I | \mathbf{P}, \mathbf{Q}, \Theta_f)$$


$$f(\mathbf{P}^T \mathbf{v}_u^U, \mathbf{Q}^T \mathbf{v}_i^I) = \phi_{out}(\phi_X(\dots \phi_2(\phi_1(\mathbf{P}^T \mathbf{v}_u^U, \mathbf{Q}^T \mathbf{v}_i^I))\dots))$$

Learning NCF

- Existing Pointwise Methods

$$L_{sq} = \sum_{(u,i) \in \mathcal{Y} \cup \mathcal{Y}^-} w_{ui} (y_{ui} - \hat{y}_{ui})^2$$

- Squared Loss explained by Gaussian Distribution (Observation)
 - But target value y_{ui} is **binarized**




- Probabilistic Approach in NCF
 - $y_{ui} = 1$ means item i is relevant to u , and 0 otherwise
 - \hat{y}_{ui} represents how likely i is relevant to u
 - Use the Logistic activation Function to make $\hat{y}_{ui} \in [0,1]$

Learning NCF

- Define Likelihood Function

$$p(\mathcal{Y}, \mathcal{Y}^- | \mathbf{P}, \mathbf{Q}, \Theta_f) = \prod_{(u,i) \in \mathcal{Y}} \hat{y}_{ui} \prod_{(u,j) \in \mathcal{Y}^-} (1 - \hat{y}_{uj})$$



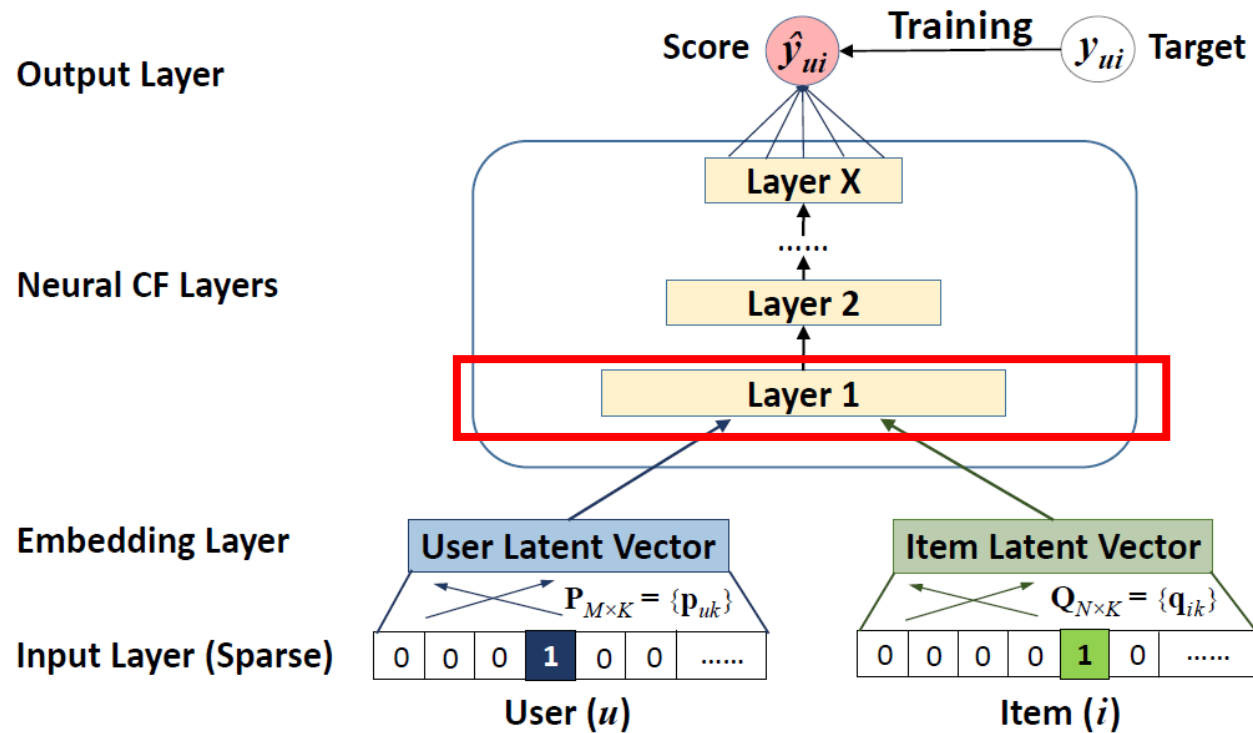
Negative Logarithm

$$L = - \sum_{(u,i) \in \mathcal{Y}} \log \hat{y}_{ui} - \sum_{(u,j) \in \mathcal{Y}^-} \log(1 - \hat{y}_{uj}) = - \sum_{(u,i) \in \mathcal{Y} \cup \mathcal{Y}^-} y_{ui} \log \hat{y}_{ui} + (1 - y_{ui}) \log(1 - \hat{y}_{ui})$$

- Same as Binary cross-entropy loss
 - Recommendation with implicit feedback as a binary classification problem
- Uniformly sample negative instances from unobserved interactions in each iteration

Generalized Matrix Factorization (GMF)

Let \mathbf{p}_u be $\mathbf{P}^T \mathbf{v}_u^U$ and \mathbf{q}_i be $\mathbf{Q}^T \mathbf{v}_i^I$



- First CF Layer

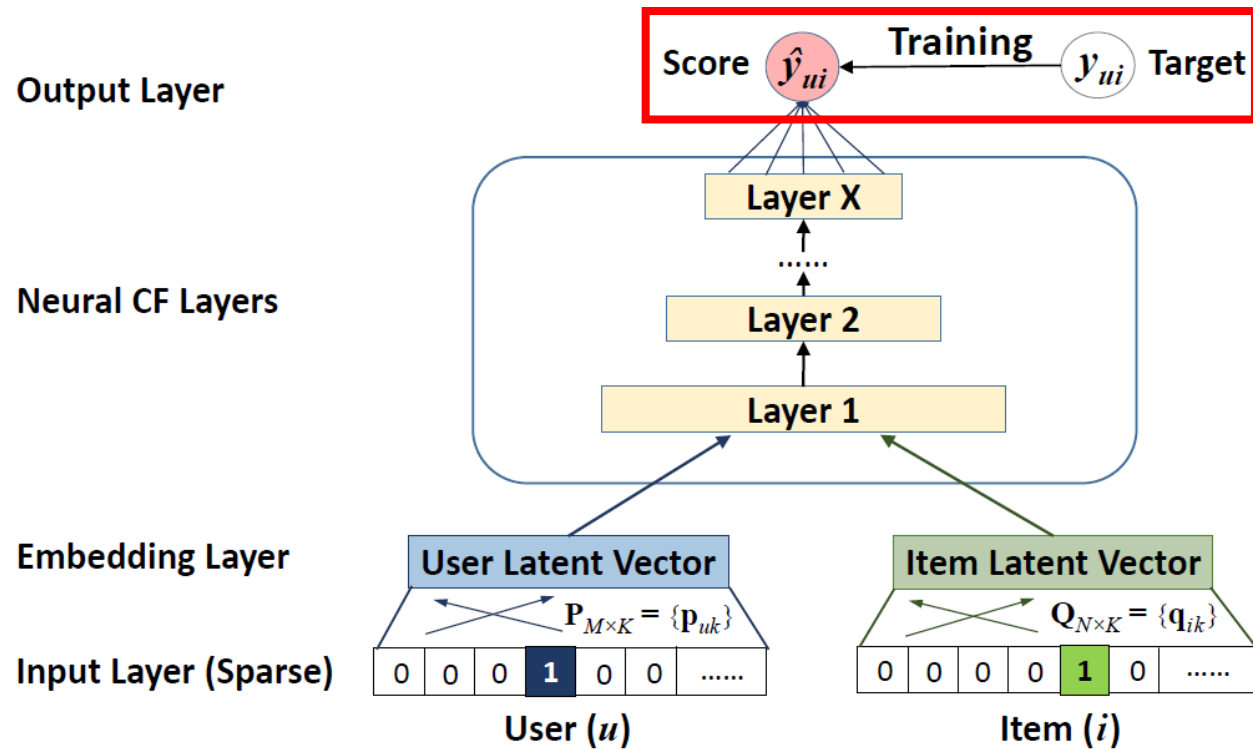
$$\phi_1(\mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u \odot \mathbf{q}_i$$

\odot denotes the element-wise product of vectors

Figure 2: Neural collaborative filtering framework

Generalized Matrix Factorization (GMF)

Let \mathbf{p}_u be $\mathbf{P}^T \mathbf{v}_u^U$ and \mathbf{q}_i be $\mathbf{Q}^T \mathbf{v}_i^I$



- Output Layer

$$\hat{y}_{ui} = \underline{a_{out}}(\underline{\mathbf{h}^T}(\mathbf{p}_u \odot \mathbf{q}_i))$$

Identity Function

Uniform Vector of 1

“Become MF”

$$\hat{y}_{ui} = \mathbf{p}_u^T \mathbf{q}_i$$

Figure 2: Neural collaborative filtering framework

Generalized Matrix Factorization (GMF)

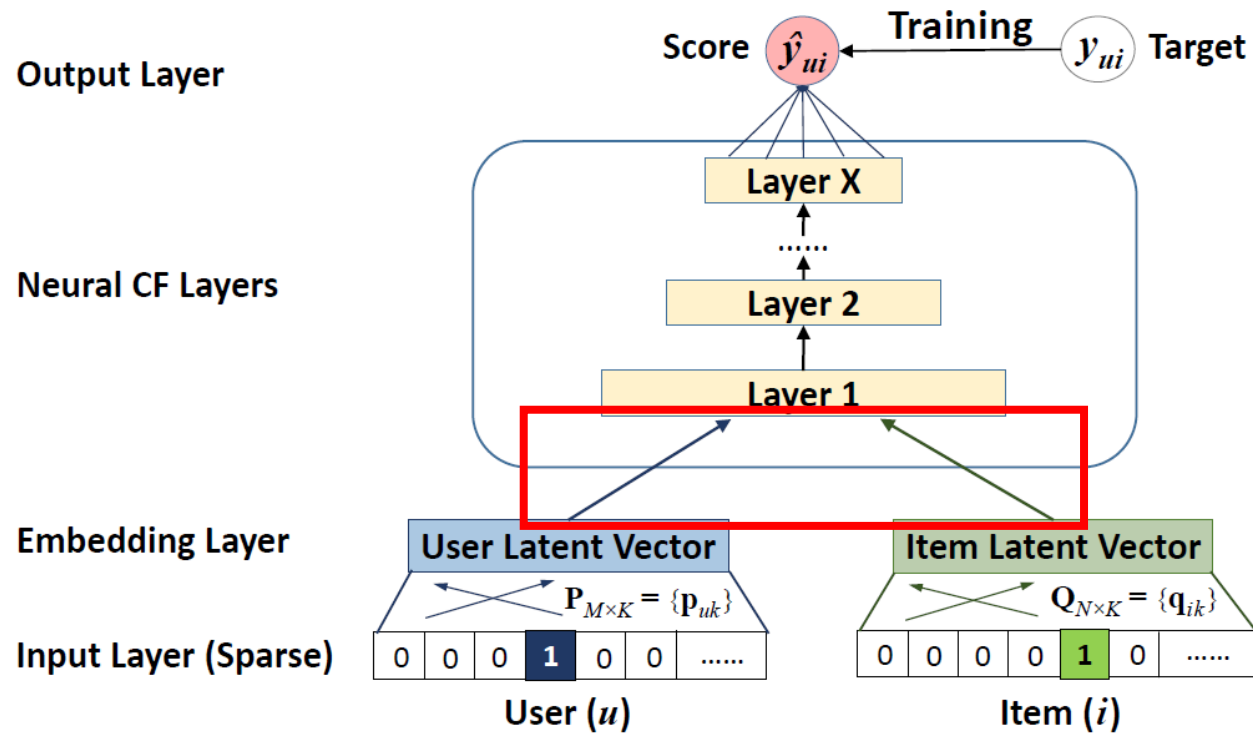
- Generalized and Extended Version of MF

$$\hat{y}_{ui} = a_{out}(\mathbf{h}^T(\mathbf{p}_u \odot \mathbf{q}_i))$$

1. Allow h to be learnt from Data
 - Vary importance of latent dimension
2. Use a non-linear function for a_{out}
 - Be more expressive than linear MF model

Multi-Layer Perceptron (MLP)

Let \mathbf{p}_u be $\mathbf{P}^T \mathbf{v}_u^U$ and \mathbf{q}_i be $\mathbf{Q}^T \mathbf{v}_i^I$



- Hidden Layers on the concatenated vector
- Layer Composition

$$\mathbf{z}_1 = \phi_1(\mathbf{p}_u, \mathbf{q}_i) = \begin{bmatrix} \mathbf{p}_u \\ \mathbf{q}_i \end{bmatrix},$$

$$\phi_2(\mathbf{z}_1) = a_2(\mathbf{W}_2^T \mathbf{z}_1 + \mathbf{b}_2),$$

.....

$$\phi_L(\mathbf{z}_{L-1}) = a_L(\mathbf{W}_L^T \mathbf{z}_{L-1} + \mathbf{b}_L),$$


$$\hat{y}_{ui} = \sigma(\mathbf{h}^T \phi_L(\mathbf{z}_{L-1})),$$

- ReLU yields slightly better

Figure 2: Neural collaborative filtering framework


Fusion of GMF and MLP

✓ GMF uses a linear kernel & MLP uses a non-linear kernel



Natural Question


? How can we fuse GMF and MLP, so that they make a better model?



1. **Share** same embedding layer and

then combine $\hat{y}_{ui} = \sigma(\mathbf{h}^T a(\mathbf{p}_u \odot \mathbf{q}_i + \mathbf{W} \begin{bmatrix} \mathbf{p}_u \\ \mathbf{q}_i \end{bmatrix} + \mathbf{b}))$

- From Neural Tensor Network
- Limit the performance
 - Same size of embeddings



✓ **Separate** embedding layer and then combine

Fusion of GMF and MLP

\mathbf{p}_u^G and \mathbf{p}_u^M denote user embeddings, \mathbf{q}_i^M and \mathbf{q}_i^G denote item embeddings

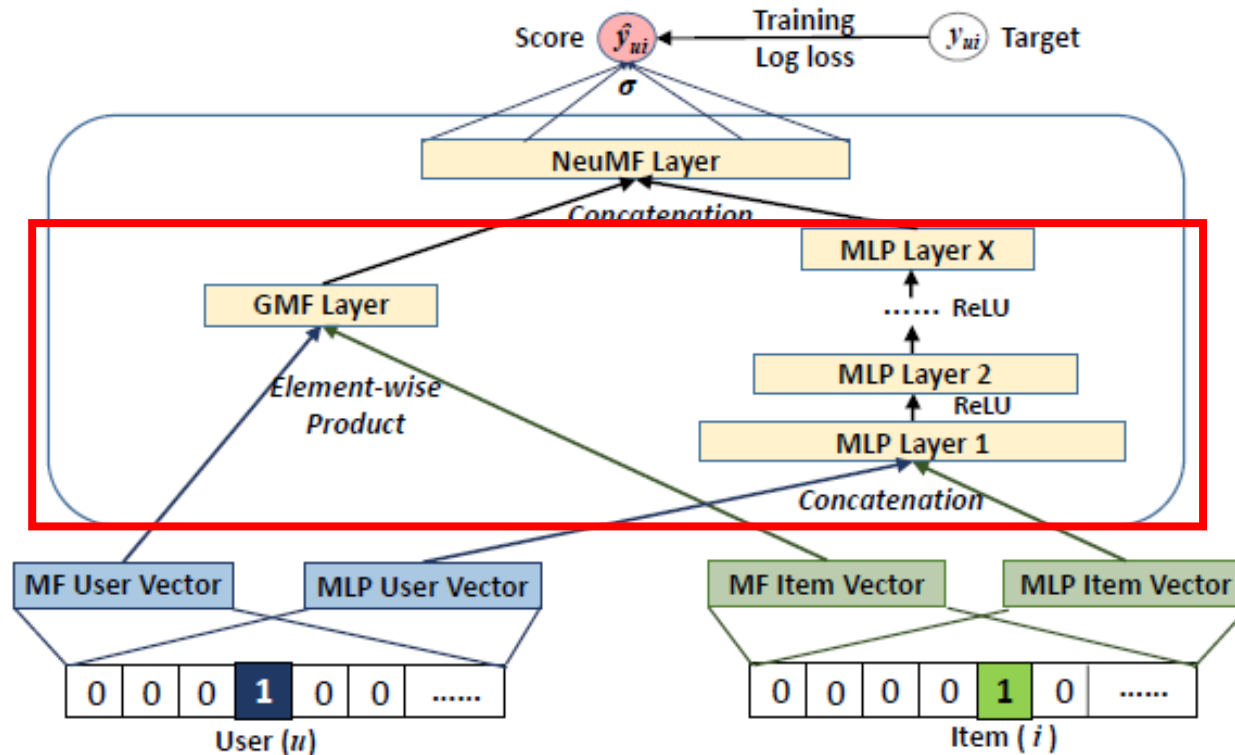


Figure 3: Neural matrix factorization model

- Combines linearity of MF and non-linearity of DNNs
- Layer Composition

$$\phi^{GMF} = \mathbf{p}_u^G \odot \mathbf{q}_i^G,$$

$$\phi^{MLP} = a_L(\mathbf{W}_L^T(a_{L-1}(\dots a_2(\mathbf{W}_2^T \begin{bmatrix} \mathbf{p}_u^M \\ \mathbf{q}_i^M \end{bmatrix} + \mathbf{b}_2)\dots)) + \mathbf{b}_L)$$

Fusion of GMF and MLP

\mathbf{p}_u^G and \mathbf{p}_u^M denote user embeddings, \mathbf{q}_i^M and \mathbf{q}_i^G denote item embeddings

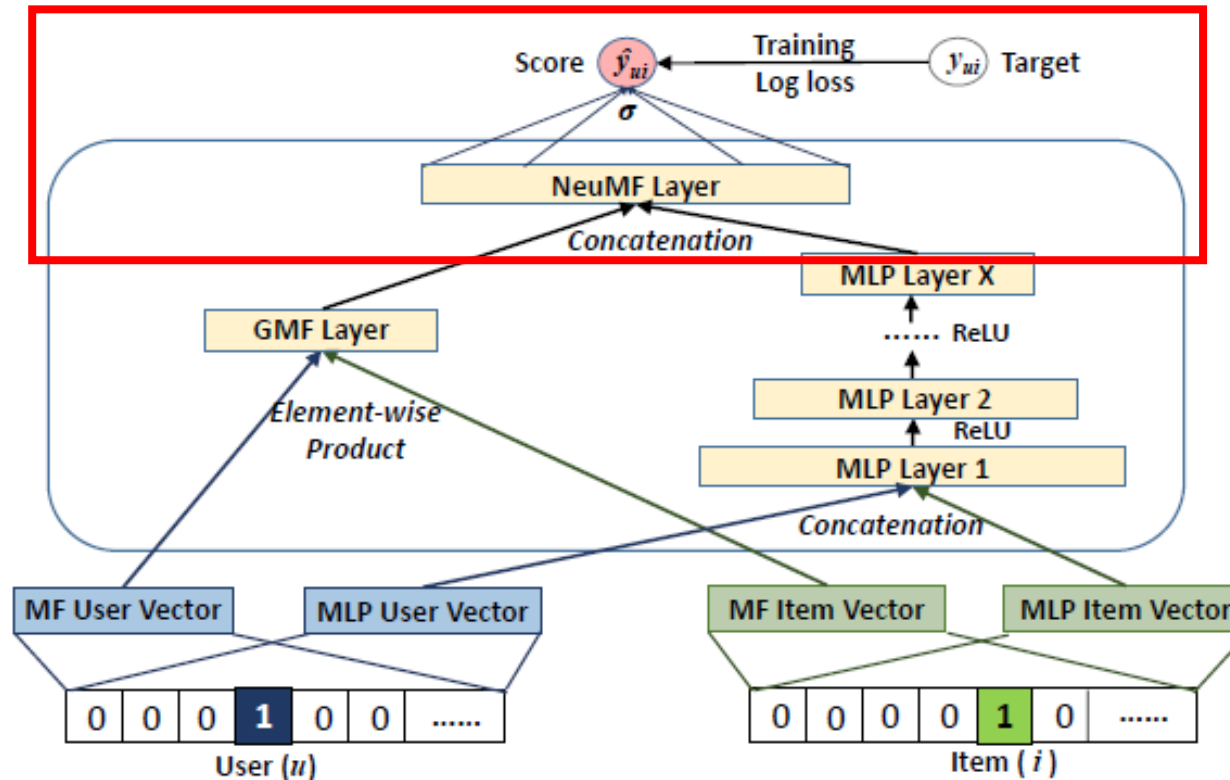


Figure 3: Neural matrix factorization model

- Hidden Layers on the concatenated vector
- Layer Composition

$$\hat{y}_{ui} = \sigma(\mathbf{h}^T \begin{bmatrix} \phi^{GMF} \\ \phi^{MLP} \end{bmatrix})$$

“Neural Matrix Factorization”

Pre-Training

- Non-convexity of the objective function of NeuMF



- Initialization determines **convergence** and **performance**



- Initialize NeuMF using the pretrained models of GMF and MLP
 1. Train GMF and MLP with random initializations
 2. Use their model parameters as the initialization

$$\mathbf{h} \leftarrow \begin{bmatrix} \alpha \mathbf{h}^{GMF} \\ (1 - \alpha) \mathbf{h}^{MLP} \end{bmatrix} \quad \bullet \text{ Weighted concatenation of pretrained } \mathbf{h} \text{ vector}$$

Research Questions

- Conduct experiments with the aim of answering three questions
 1. Do NCF methods **outperform**?
 2. **How does** optimization framework **work** for the recommendation task?
 3. Are **deeper layers** of hidden units **helpful** for learning from data?

Experimental Settings

I. Datasets

1. MovieLens

- Movie rating dataset
 - Learning from implicit signal of explicit feedback
- Whether user has rated item

2. Pinterest

- SNS for sharing images
 - Constructed by below paper
 - X. Geng, H. Zhang, J. Bian, and T.-S. Chua. Learning image and user features for recommendation in social networks. In ICCV, pages 4274–4282, 2015.

Table 1: Statistics of the evaluation datasets.

Dataset	Interaction#	Item#	User#	Sparsity
MovieLens	1,000,209	3,706	6,040	95.53%
Pinterest	1,500,809	9,916	55,187	99.73%

Experimental Settings

II. Evaluations Protocols

- Leave-one-out (latest interaction as test, remaining for train)
- Rank the test item among 100 unobserved randomly sampled items

1. HR (Hit Ratio)

- Whether test item is present on the top-K list

2. NDCG

- The position of the hit by assigning higher scores to hits at top ranks

III. Baselines

1. ItemPop

- Non-personalized

2. ItemKNN

- Item-based CF

3. BPR

- Pairwise Loss

4. eALS

- State-of-the-art MF

Experimental Settings

IV. Parameter Settings

- To determine hyperparameters, randomly sampled one interaction for each user as the validation data
- Initialized model parameters with a Gaussian Distribution – $N(0, 0.01^2)$
- Optimizing with mini-batch Adam (batch size – [128, 256, 512, 1024])
- Learning rate – [0.0001, 0.0005, 0.001, 0.005]
- Last hidden layer determines the model capability – term it as predictive factors [8, 16, 32, 64]
 - Ex) Architecture of NCF layers 32 → 16 (embedding size) → 8, three hidden layers
- Set alpha = 0.5 for pre-training

4. Experiments

Performance Comparison (RQ1)

- Performance Result

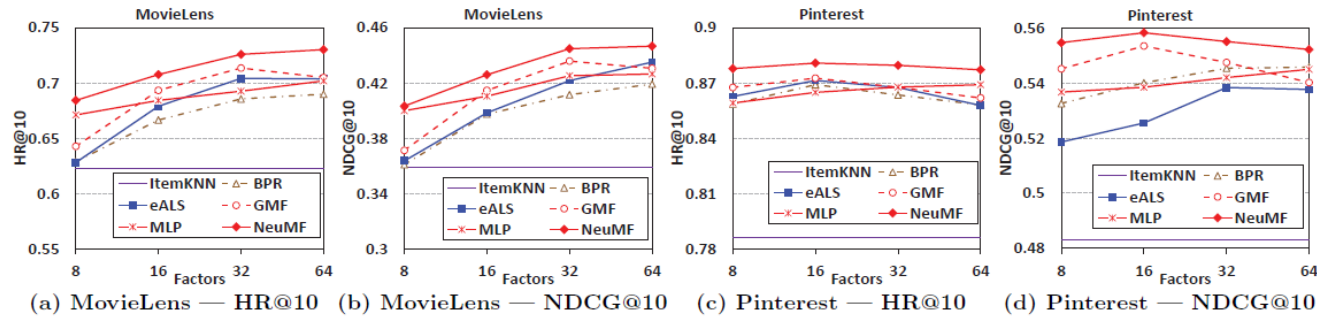


Figure 4: Performance of HR@10 and NDCG@10 *w.r.t.* the number of predictive factors on the two datasets.

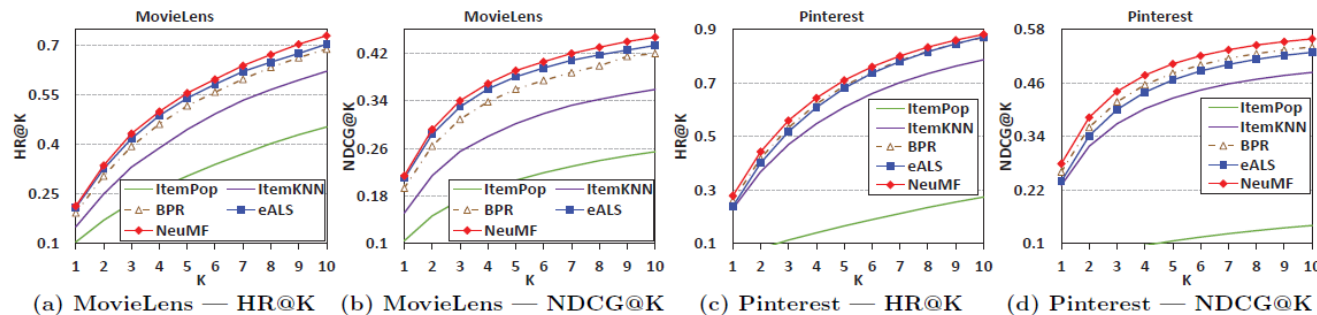


Figure 5: Evaluation of Top-K item recommendation where K ranges from 1 to 10 on the two datasets.

- ④ NeuMF achieves the **best performance** on both datasets
- ④ GMF and MLP also show **quite strong performance**
- ④ GMF shows consistent improvements **over BPR**
- ⑤ NeuMF demonstrates **consistent improvements**

Performance Comparison (RQ1)

- Utility of Pre-training

Table 2: Performance of NeuMF with and without pre-training.

Factors	With Pre-training		Without Pre-training	
	HR@10	NDCG@10	HR@10	NDCG@10
MovieLens				
8	0.684	0.403	0.688	0.410
16	0.707	0.426	0.696	0.420
32	0.726	0.445	0.701	0.425
64	0.730	0.447	0.705	0.426
Pinterest				
8	0.878	0.555	0.869	0.546
16	0.880	0.558	0.871	0.547
32	0.879	0.555	0.870	0.549
64	0.877	0.552	0.872	0.551

- NeuMF with pre-training achieves better performance in most cases

4. Experiments

Log Loss with Negative Sampling (RQ2)

- Result of Pointwise Log loss




Figure 6: Training loss and recommendation performance of NCF methods *w.r.t.* the number of iterations on MovieLens (factors=8).

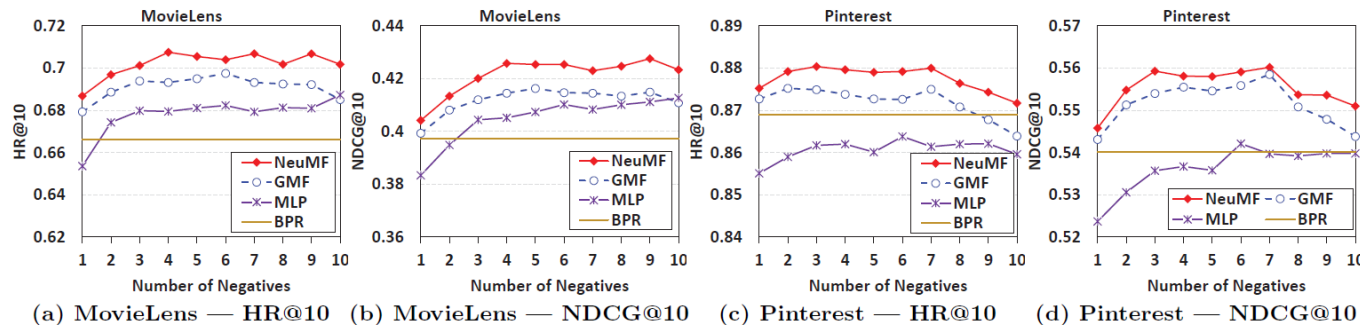


Figure 7: Performance of NCF methods *w.r.t.* the number of negative samples per positive instance (factors=16). The performance of BPR is also shown, which samples only one negative instance to pair with a positive instance for learning.

⑥ With more iterations, loss gradually decreases

⑥ NeuMF achieves the lowest loss followed by MLP, GMF

⑦ Pointwise loss can handle flexible sampling ratio

⑦ Too large sampling ratio hurts the performance

Is Deep Learning Helpful? (RQ3)

- MLP with different layers

Table 3: HR@10 of MLP with different layers.

Factors	MLP-0	MLP-1	MLP-2	MLP-3	MLP-4
MovieLens					
8	0.452	0.628	0.655	0.671	0.678
16	0.454	0.663	0.674	0.684	0.690
32	0.453	0.682	0.687	0.692	0.699
64	0.453	0.687	0.696	0.702	0.707
Pinterest					
8	0.275	0.848	0.855	0.859	0.862
16	0.274	0.855	0.861	0.865	0.867
32	0.273	0.861	0.863	0.868	0.867
64	0.274	0.864	0.867	0.869	0.873

Table 4: NDCG@10 of MLP with different layers.

Factors	MLP-0	MLP-1	MLP-2	MLP-3	MLP-4
MovieLens					
8	0.253	0.359	0.383	0.399	0.406
16	0.252	0.391	0.402	0.410	0.415
32	0.252	0.406	0.410	0.425	0.423
64	0.251	0.409	0.417	0.426	0.432
Pinterest					
8	0.141	0.526	0.534	0.536	0.539
16	0.141	0.532	0.536	0.538	0.544
32	0.142	0.537	0.538	0.542	0.546
64	0.141	0.538	0.542	0.545	0.550

- Stacking more layers are beneficial to performance
 - High non-linearities brought by stacking more non-linear layers
- Simply concatenating latent vectors is insufficient for modelling (MLP with no hidden layers)

Conclusion and Future Work

- Explored Neural Network Architectures for CF
 - Simple and general framework NCF
 - Three instantiations – GMF, MLP, NeuMF
 - Opening up a new avenue for recommendation based on Deep Learning
- Future Work
 - Study pairwise learners for NCF
 - Extend NCF to model auxiliary information
 - Such as user review, knowledge bases, temporal signals
 - Build recommender systems for multi-media items
 - Effective methods for multi-view and multi-modal data

Python Implementation

- To be updated (Not complete yet)
 - Having trouble dealing with argparse and non-csv file
 - <https://github.com/LeeJunmo/DSAIL-Lab-Intern>

Any Questions?